


Learning to Align Speech-Text with Optimal Transport



t-SNE visualization of encoder outputs

Color codes for embeddings:

Text

Speech (cross-entropy loss)

Speech (CTC loss)

Speech (CTC + OT loss)

Model trained with optimal transport produces well aligned speech-text features.